

Apache Spark™

Lightning-Fast Cluster Computing

Overview and use cases

Me

- ▶ “Hadoopified” and deeply fell in love with “Big Data” in 2010
- ▶ Software development Engineer – Relational DBs – @Amadeus – (2 years)
- ▶ “Big” Data Engineer – Hadoop , Data Analysis & NoSQL @Amadeus – (1 year)
- ▶ Now starting to be “Sparkified” ;–)



Agenda

- ▶ Big Data
 - Short overview
 - New challenges
- ▶ Hadoop – the Big elephant that solves “Big” problems
 - Overview
 - Hadoop map-reduce – Advantages and limitations
- ▶ Spark – will make the Big elephant “Fly”
 - Short overview
 - Architecture
 - Demo-use cases : bookings analysis, flight delay prediction (Spark SQL, Spark Streaming, Machine Learning)

1 — Big Data

Big Data – short explanation

A graceful mixture of all three V : (Big) **Variety**, (Big) **Velocity**, (Big) **Volume** of Data



2000 – 2010
Enterprise Data Lake



2010 –
Continuous stream of data (structured and unstructured)

1980 – 2000
Enterprise data warehouses



Big Data – new challenges

- ▶ A single machine can't deal with Big – Volume, Variety, Velocity of Data
 - Solution : use many machines (*sounds trivial* 😊)



- ▶ New Challenges
 - Distributed Storage
 - Distributed Processing
 - Communication between machines
 - Network and machine failure
 - ...



2

Hey we found a solution !



Hadoop

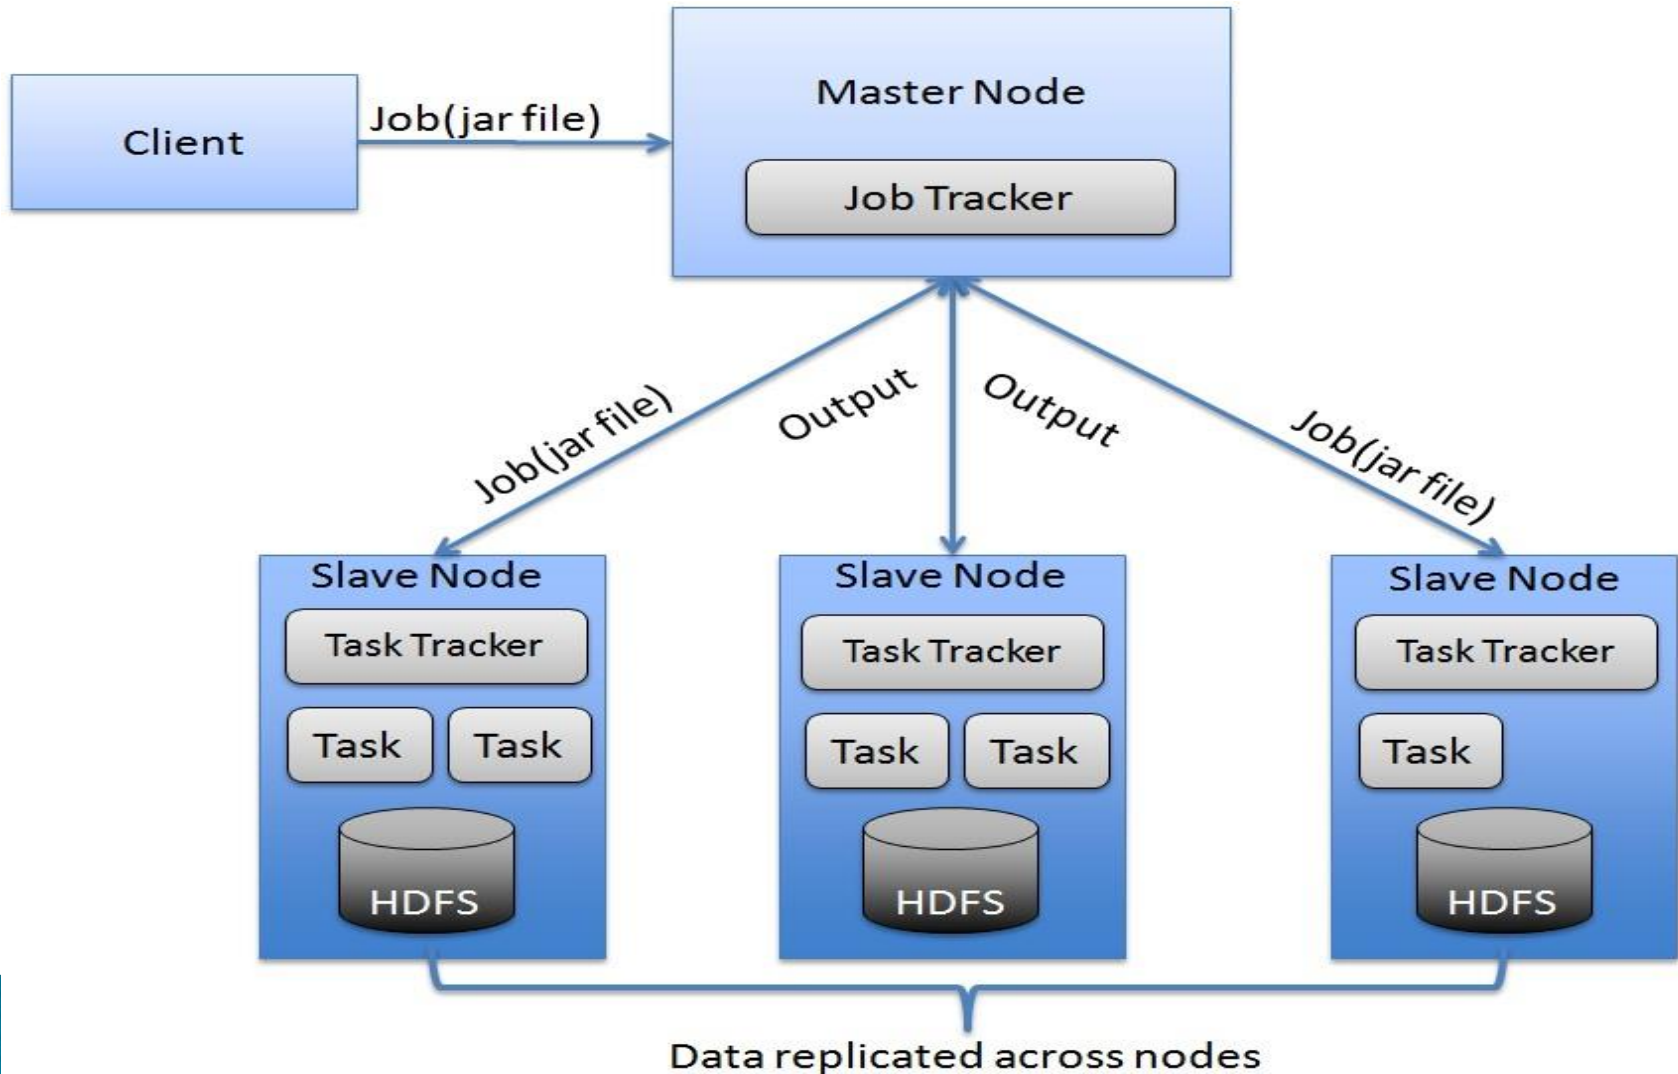
Big Data Platform

▶ Hadoop platform

- Distributed file system (hdfs)
- Parallel computing framework (originaly Map/Reduce)
- Fault tolerant, handle replication, Data locality , node failures, ...
- Scalable (Yahoo! has a 4000–node cluster)
- High performances : sorted a TB of random integers in 62 seconds
- Cheap 😊 : deploy on commodity machines

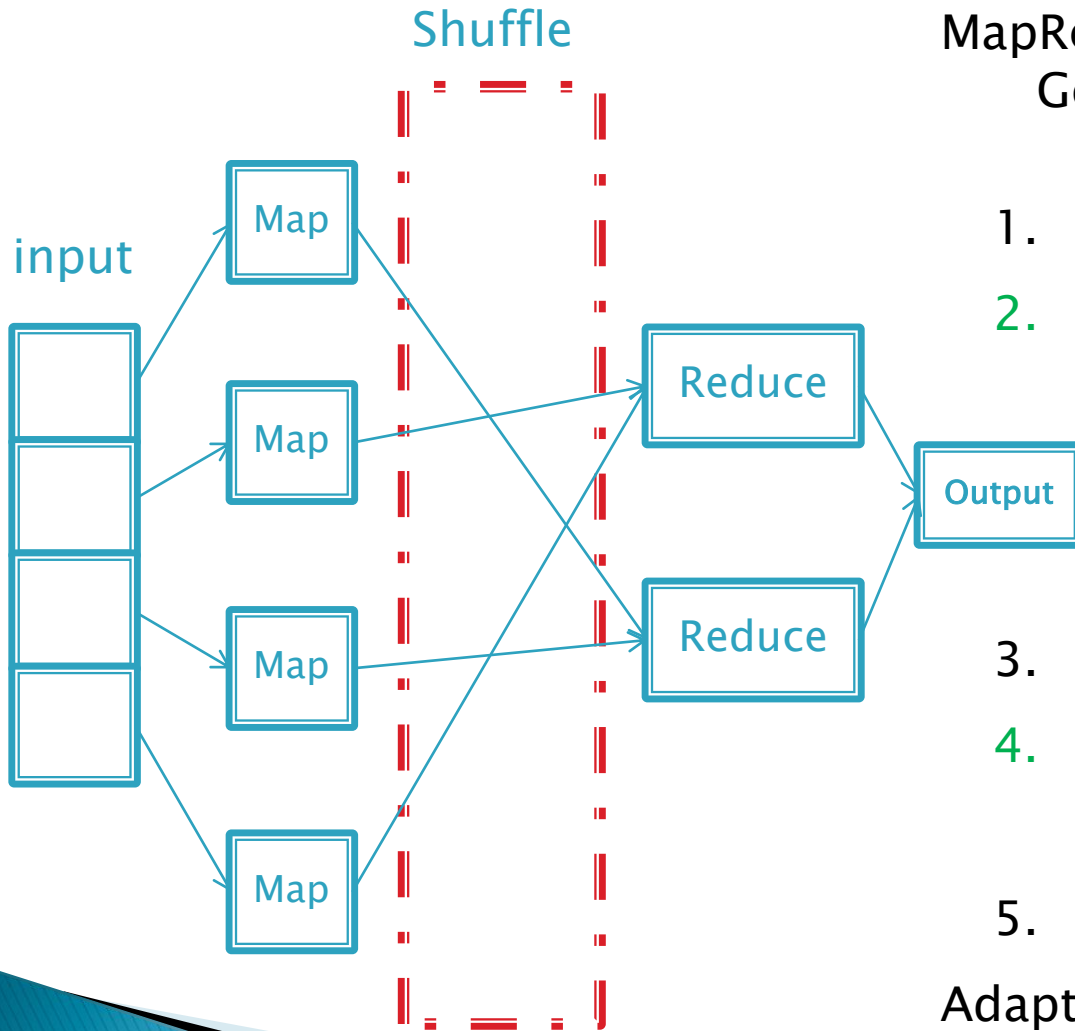
Hadoop

Architecture



Hadoop

Map-Reduce : core processing framework



MapReduce Framework
Google patent (2004)

1. Read a lot of data
2. **Map**: extract something you care about
3. Shuffle and sort
4. **Reduce**: aggregate, summarize, filter, transform
5. Write the results


Adapt Map and Reduce part to solve different use cases

Hadoop

Map-Reduce limitations

- ▶ Relies on *persistent storage (disk)* to provide *fault tolerance* between **map**, **shuffle** and **reduce** phases
 - As a consequence persistent storage introduces higher latency
- ▶ It's one-pass computation model makes it a *poor fit* for *low-latency applications* and *iterative computations*, such as *machine learning and graph algorithms*.
- ▶ Hadoop Map-Reduce **not good** at:
 - More **complex**, multi-stage applications (graph algorithms, machine learning)
 - More **interactive** ad-hoc queries
 - More **real-time** online processing
- ▶ All three of these apps require *fast data sharing* across parallel jobs.

3 — Well... let's make the elephant **Fly** !




Apache Spark

overview

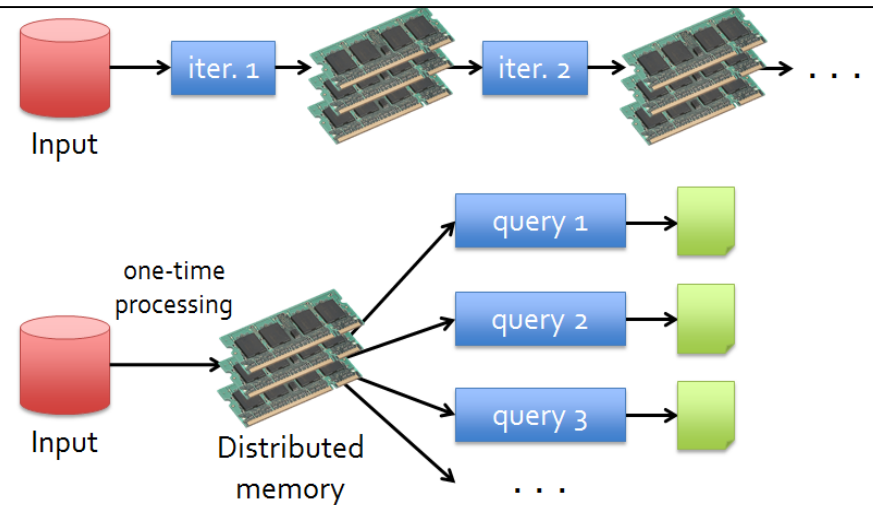
- ▶ fast and general engine (map-reduce like but not restricted) for large-scale data processing
- ▶ *In-memory* (as much as possible) data storage for very fast iterative queries, with *fault-tolerance* mechanisms
- ▶ Compatible with Hadoop's storage APIs

Hadoop Map-Reduce



Slow due to replication, serialization, and disk IO

Apache Spark



10-100x faster than network and disk

Spark Main Goal

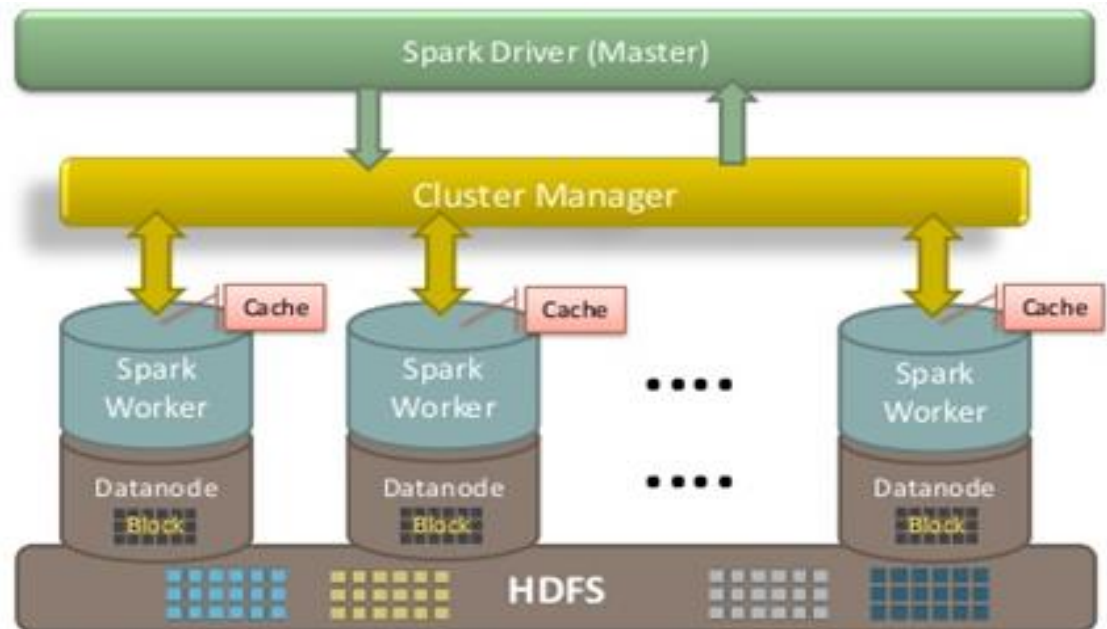
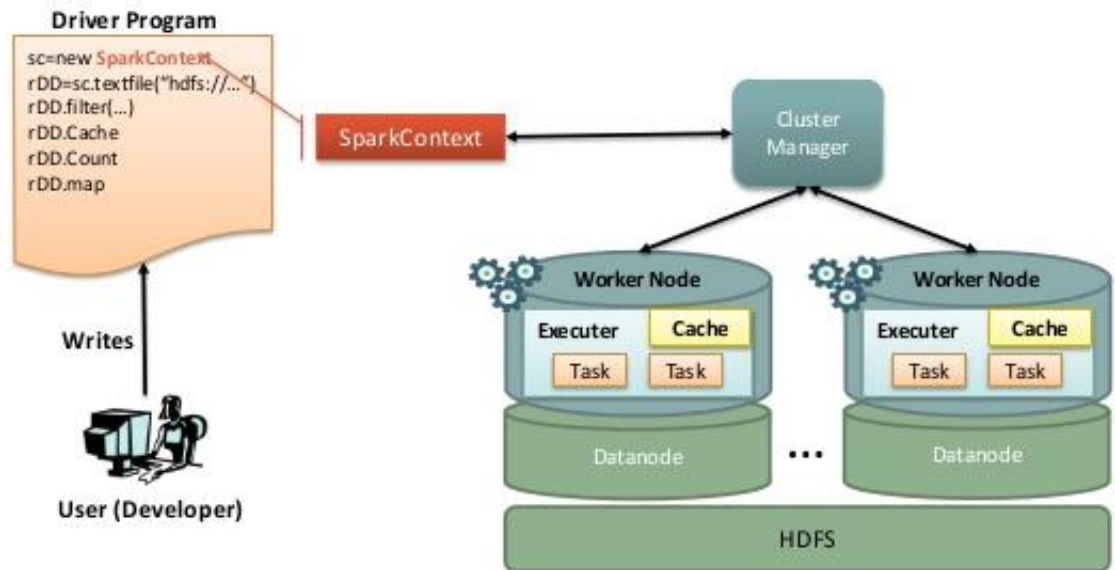
- ▶ Provide distributed *memory* abstractions for clusters to support apps with working sets
- ▶ Retain the attractive properties of Hadoop MapReduce:
 - Fault tolerance (for crashes & stragglers)
 - Data locality
 - Scalability

Challenge : How to design a distributed memory abstraction that is both fault-tolerant and efficient ?

Solution : augment data flow model with “resilient distributed datasets” (*RDDs*)


Spark

Architecture



Credit :
www.unicomlearning.com

Spark Goal: unification



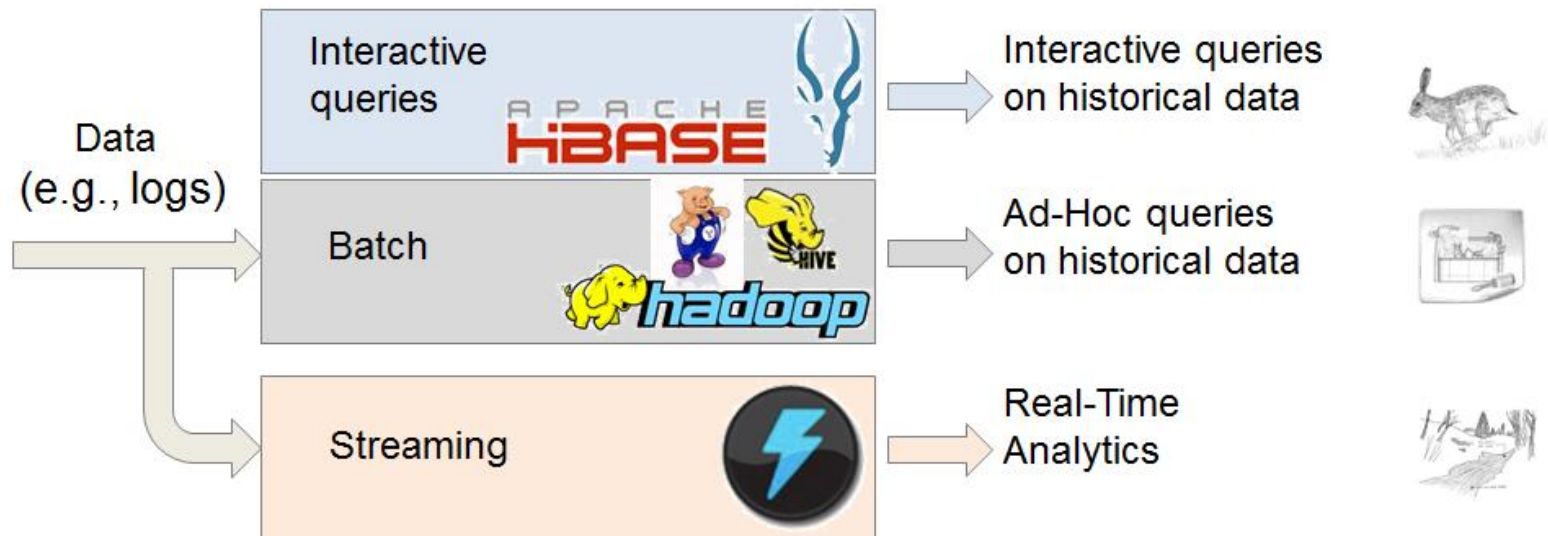
Support **batch**, **streaming**, and **interactive** computations...

... in a **unified** framework... allowing **code-reuse** across all type of computations

Easy to develop **sophisticated** algorithms (e.g., graph processing , machine learning, near real-time analytics)

The Need For Unification

► Today's state-of-art analytics stack




Challenge 1 : need to maintain three stacks

- Expensive and complex
- Hard to compute consistent metrics across stacks

Challenge 2 : hard/slow to share data, e.g.,

- » Hard to perform interactive queries on streamed data



What's the magic behind Spark ?

Spark RDD

RDD (*Resilient distributed dataset*) :


- *distributed memory abstraction*
- enables to perform *parallel in-memory* computations on large clusters in a *fault-tolerant* manner.
- can be considered as a *read-only, partitioned collection of records*.
- can be re-computed easily thanks to it's list of parent RDDs

RDD *interface* :


- A set of partitions
- A set of Dependencies on parent RDDs
- A function (compute) for computing it from its parents
- Optionally a partitioner for key-value RDD (e.g to say that the RDD is hash partitioned)
- Optionally a list of preferred locations to compute each split

Example Spark code snippet

```
val lines = spark.textFile("hdfs://.../serverlogs")  
val errors = lines.filter(_.startsWith("ERROR"))  
val tsField = errors.map(_.split('\t')(3))
```



3 — Ok! Now let's see all this in **Play** !



Demo1_____

Flights Bookings count per (departure, destination) airport pair – in Batch

Bookings log data

BookingID, BookingTimestamp , CarrierCode, FlightID, CustomerID, AgeCategory ,
DepartureAirport, **DestinationAirport**, DepartureDate, ReturnDate, NbPassengers ,
TicketPrice

1,20141204-103000,AF,153,2145,SENIOR,**NCE,CDG**,20141220,20141231,2,150
2,20141204-103005,AF,124,1462,ADULT,**LYS,MRS**,20141215,20150105,1,115
3,20141206-154500,BA,326,4527,ADULT,**LGW,BOD**,20150107,20150121,1,200
4,20141207-112000,AL,147,5982,INFANT,**NCE,CDG**,20150204,20150210,1,100
....



(NCE,CDG) , 2
(LYS,MRS) , 1
(LGW,BOD), 1


Demo2_____

Flights Bookings count per (departure, destination)
airport pair – Streaming (near real-time)

Booking server



logs



users

Booking logs generated in real-time

```
BookingID, BookingTimestamp .... , DepartureAirport,  
DestinationAirport, ....  
1,20141204-103000, AF,153,2145,SENIOR,NCE,CDG,  
3 20141204-103000, LF,147,5982,INFANT,NCE,CDG, ...  
2,20141204-103002, AF,124,1462,ADULT,LYS,MRS,201...  
....
```



+ **Spark** SQL

```
Time : 20141204-103000  
-----  
(NCE,NICE,CDG,PARIS) , 2  
Time : 20141204-103002  
-----  
(LYS,LYON,MRS,MARSEILLE) , 1
```


Demo3_____

Flights Delay prediction – Machine Learning

Training set

Year,Month,...FlightNum,
DepartureDelay,...
2007,1,...,256,17
2007,3,...,365,10

Test set

Year,Month,...FlightNum,
DepartureDelay,...
2007,6,...,358,12
2007,8,...,627,14

Live flights data

Year,Month,...FlightNum,
DepartureDelay,...
2008,7,...,256,17
2008,11,...,365,10

**Spark
MLib**
(machine
learning)

**Flight
Delays
Prediction
Model**


**Spark
Streaming**

**Spark
MLib**
(machine
learning)

Flight delays prediction on live data

Flight No 7, from ORD, destination = HLN:prediction for delay= 1.0 , actual delay=1.0
Flight No 11, from ORD, destination = PHX:prediction for delay= 1.0, actual delay=1.0
Flight No 1, from DEN, destination = GEG:prediction for delay= 1.0, actual delay=0.0
...

Sparkling Thanks !



Questions ?

